


Convolutional Neural Networks


INTRODUCTION TO IMAGE CLASSIFICATION
(MODULE 6 – VISION MODELS) – By Christelle JULIAS



Content


- ❑ Discovery and Overview
- ❑ [Computer Vision - Images]
- ❑ Use Cases of CNNs IRL
 - ❑ Image Classification
 - ❑ Object Recognition
- ❑ What are Convolutional Neural Networks?
- ❑ How do they work?
 - ❑ Core Components: Convolution, Filters, and Pooling layers
- ❑ Let's explore!
- ❑ How do we manipulate Convolutional Neural Networks?
 - ❑ Creation
 - ❑ Training
 - ❑ Fine-tuning
 - ❑ A bit of Maths

Muffins or Chihuahuas?




Introduction

WHAT CAN WE DO WITH IMAGES?



Images are « sets of pixels »

1. A **Pixel** (short for "picture element") : smallest unit of a digital image or display.
2. Represents a **single point** in an image. Measure of size for calibration. Larger pixels capture more light, smaller provide finer details.
3. Displays various colors based on its **RGB** (red, green, blue) **components**. In Black & White: binary, in Grayscale: B, W & shades of gray.



- **Composition:** made up of subpixels, typically three for RGB color representation.
- **Resolution:** total number of pixels in an image. Ex: a 1920 x 1080 display has over 2 million pixels.
- **Function:** More pixels generally lead to clearer and more detailed images=> impact the overall quality of visual displays (ex.FullHD)


Image Processing techniques include

- **Enhancement:** parameters such as brightness, contrast, and sharpness
- **Restoration:** revamp degraded original images
- **Segmentation:** divide into "semantic/meaningful" parts according to the pixel distribution.




- **Compression:** to reduce the file size through lossy or lossless methods (MPEG, HEIF)
- **Generation:** GANs (Generative Adversarial Networks) to create new images or enhance
- **Morphological Processing:** shape image structures (twist, wrinkle etc.)

Examples



Semantic Segmentation

Ex: spatial data within contexts (streets, hours, references with self-driving cars), pixel level classification



Restorative Filters

Ex: Antique Photos (here Victor Hugo, the French Novelist of the 19th Century, by Bertall 1867)



Morphological Processing + Generation

Ex: Beauty filters (Instagram, Adobe, PicsArt ...), Background removal and replacement

Types of Object Recognition Algorithms



Image classification

- **-Classifies** / labels the object contained in the picture (what is it?)
- Provides a **probability** of the guess~(it's 90% likely to be a bird)


Traditional CNNs



Both Classification & Localization

- **- Detects** the objects and also **localizes** it in the picture simultaneously (where is the bird?)

Simplified Yolo, R- CNNs
(You Only Look Once, Region CNN models)



Object Detection

- **- Detects** the content of an image, presence of **several objects** (many bounding boxes) (what are those?)

Yolo, R- CNNs

QUIZZ



Identify real world
application of Image
Classification



Image Classification / Recognition


[Is it a [label, name]?]

- Classifying images into predefined **categories**.
- Recognizing **patterns** and **features**.
- Widely used in applications like: **photo tagging** in social media and **medical imaging** for diagnostics identifying pathologies.



```
[64]: def plot(image_batch, label_batch):  
    plt.figure(figsize=(10,5))  
    for i in range(10):  
        ax = plt.subplot(2,5,i+1)  
        img = cv2.imread(str(image_batch[i]))  
        img = cv2.resize(img, (224,224))  
        plt.imshow(img)  
        plt.title(label_batch[i])  
        plt.axis("off")
```

```
[65]: plot(image_batch, label_batch)
```



Object Detection

(Presence: Is there an [object name] in the image / landscape/ supermarket aisle etc.?)

➔ - **Identifying** and **locating** objects within a picture simultaneously.

➔ - Applications include **self-driving cars, facial recognition, security cameras, sorting/recycling wastes, tracking wildlife** etc






Image: ywseo



Face Recognition

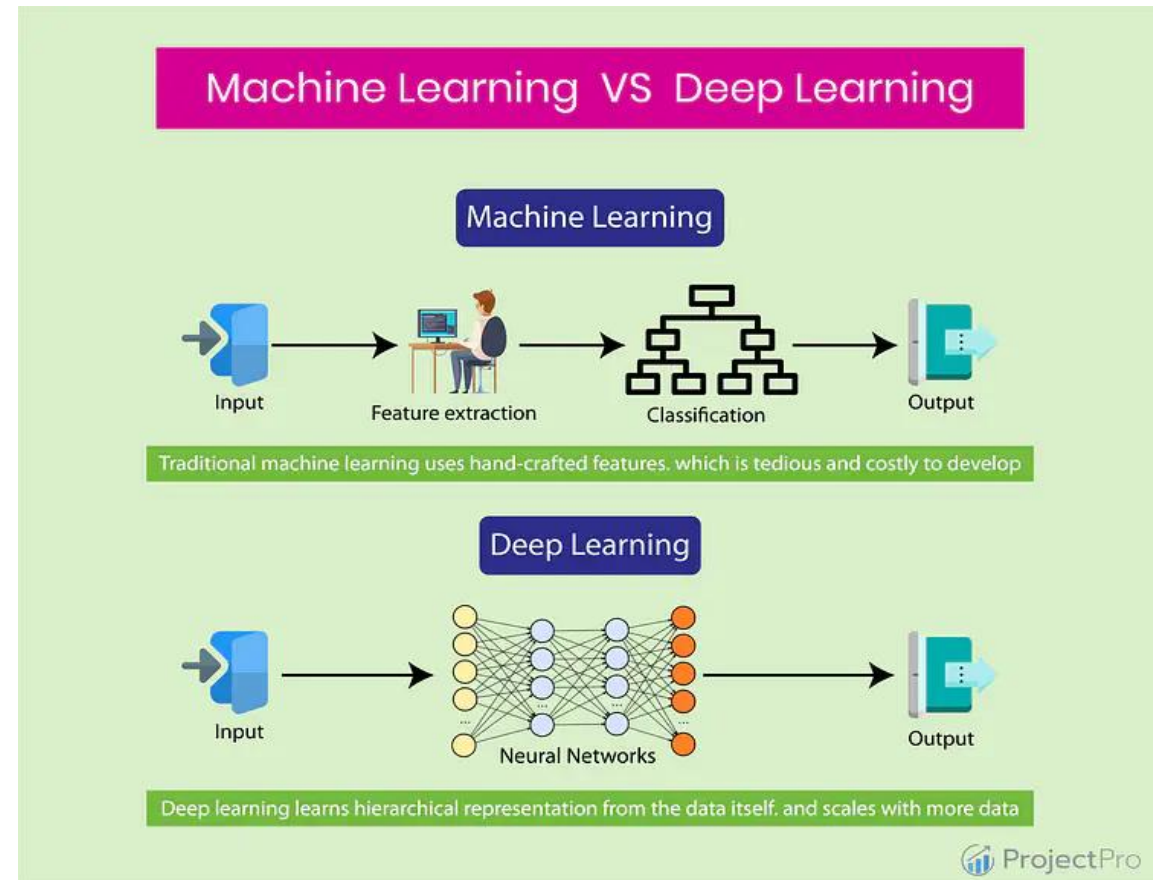
[is it {person/celebrity_name}/?]

- Attaching clues or “landmarks”, to distinctive features
- **Biometric screening**
- **ID tracking** across video frames
- **Emotional intelligence** (microexpressions)



Machine Learning vs Deep Learning

- ▶ Traditional **machine learning** uses hand-crafted features which is tedious and costly to develop
- ▶ **Deep Learning** learns hierarchical representation from the data itself and scales with more data
- ▶ Deep Learning is a subset of Machine Learning. They are complementary.



Deep Learning approach #1


Supervised Learning

Requires **labeled datasets** for training models (e.g., convolutional neural network for image classification)

How Supervised Machine Learning Works


STEP 1

Provide the machine learning algorithm categorized or "labeled" input and output data from to learn



STEP 2

Feed the machine new, unlabeled information to see if it tags new data appropriately. If not, continue refining the algorithm




TYPES OF PROBLEMS TO WHICH IT'S SUITED



CLASSIFICATION

Sorting items into categories



REGRESSION

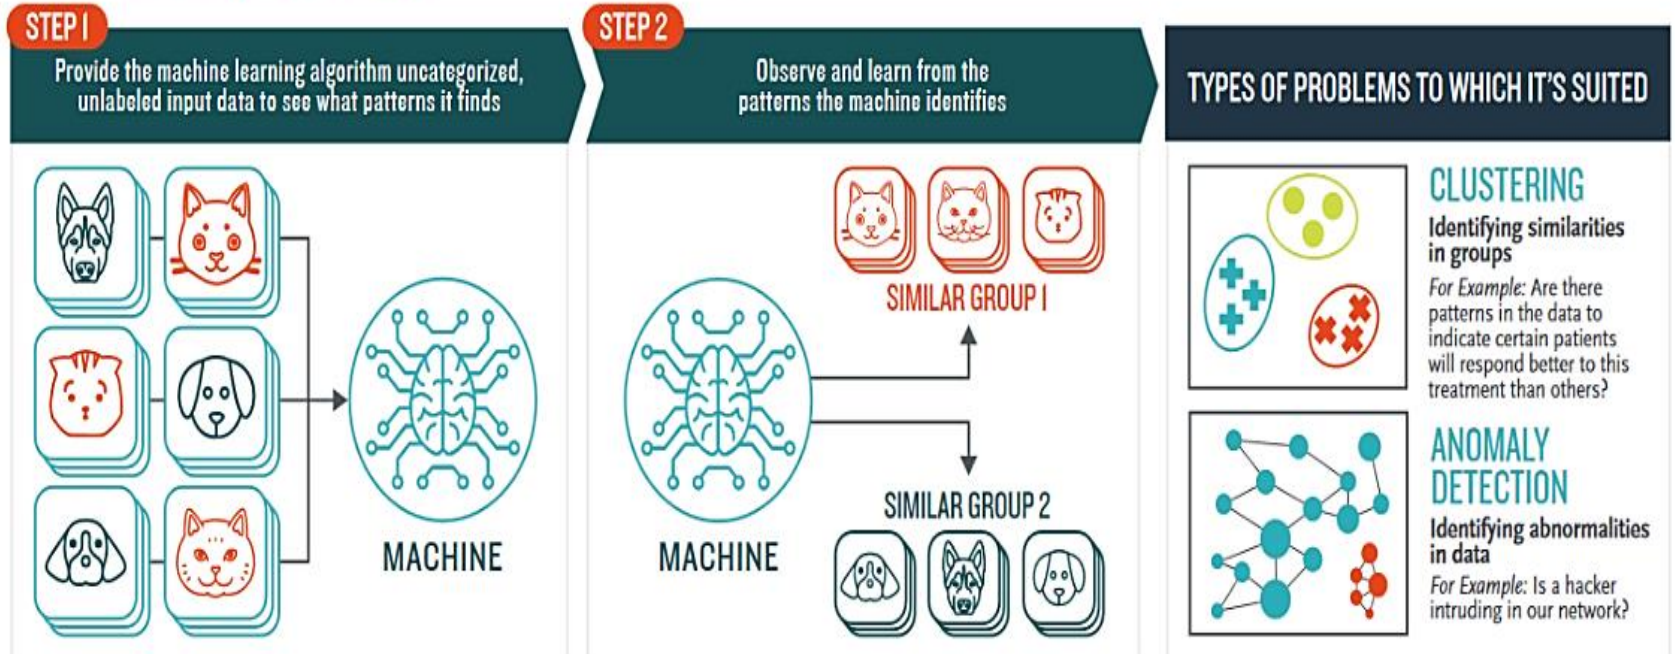
Identifying real values (dollars, weight, etc.)

Deep Learning approach #2

Unsupervised Learning

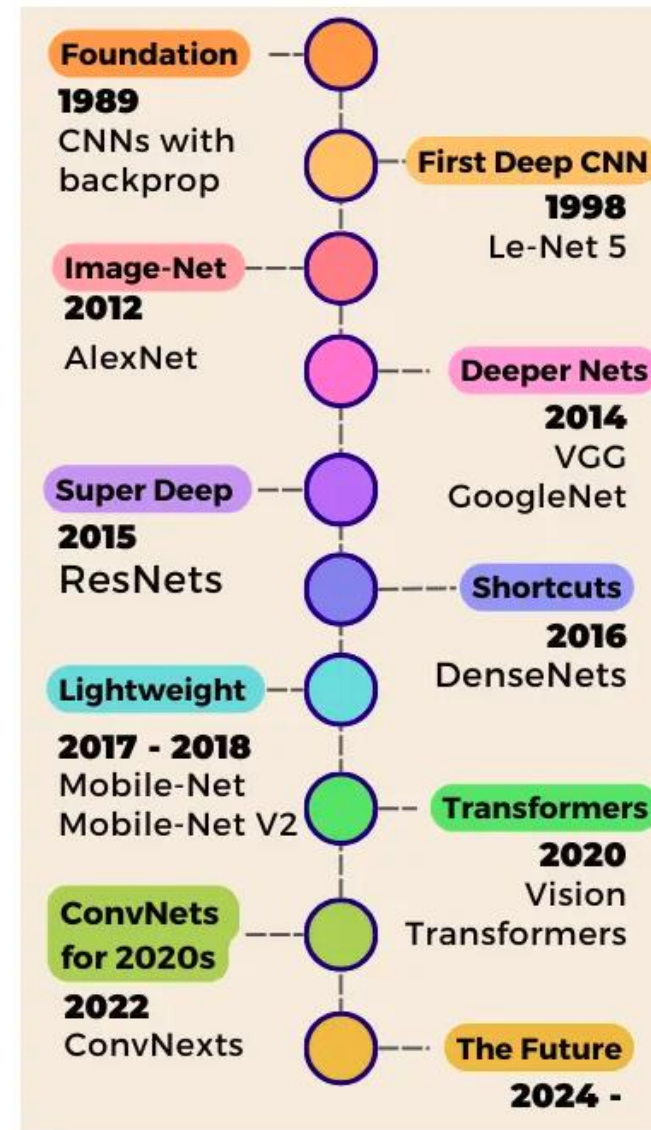
utilizes methods, such as autoencoders and generative adversarial networks (GANs) to learn from unlabeled data.

How **Unsupervised** Machine Learning Works



What are Convolutional Neural Networks?

CONVNETS OR CNNS




In the context of mathematics



$$(f * g)(t) \stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau$$

To convolve a kernel with an input signal:
flip the signal, move to the desired time,
and accumulate every interaction with the kernel

Example functions




Convolution operation result



In the context of AI, what is a Convolutional Neural Network?


1. A type of **Deep Learning model** designed for **image processing**, also called ConvNet or CNN
2. **Mimics visual perception** in **humans**
3. A powerful tool for **visual recognition** and **style manipulation**
4. Consists of multiple **sequential layers** (convolutions, filters, and pooling) to **extract features** (edges, texture) from images ex: the cat has pointy ears, and a round face
5. Essential for modern applications in **Computer Vision** (multimodal multimedia: video, images)



Architecture of a CNN

→ Deep CNN


Deep Convolutional
Network **(DCN)**



Index

- Backfed Input Cell
- Input Cell
- Noisy Input Cell
- Hidden Cell
- Probablisticc Hidden Cell
- Spiking Hidden Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- Different Memory Cell
- Kernel
- Convolutional or Pool

CNN for image classification




The Deer class gets 94% of accuracy



Image Classification with CNN

Classification is essentially a guess deduced from the extracted features and object detection may follow it once the object has been recognized.

- 1- **Feature extraction** performed by CNNs
- 2- **Pooling** as a dimensionality reduction of feature maps
- 3- **Classification**, once the extracted features pass through fully connected layers (often using Softmax)
- 4- **Object Detection**: by predicting bounding boxes around detected objects




Convolutions



1D AND 3D




What are Convolutions?

- ➔ - Definition: A mathematical operation that combines two functions.
- Purpose: To extract high-level features from images.
- Process: Apply a filter (kernel) over the image, computing dot products.
- ➔ -



Pooling Layers?


- ➔ **Function?** Reduces **spatial dimensions** of the feature maps.
- ➔ **Why?** Images can be voluminous (resolutions, size)
- ➔ **Solution:** Focus on the essential features
- ➔ **How?** Sliding a window and taking only one value
 - Types:
 - **Max Pooling:** Takes the maximum value from a set of values.
 - **Average Pooling:** Takes the average value.
 - **Benefits:** Decreases computational load and helps prevent overfitting.



Filters / Kernels

Small matrices used in convolutions.

- Detect specific features (Edges, textures).
- Different filters highlight different aspects of the image.

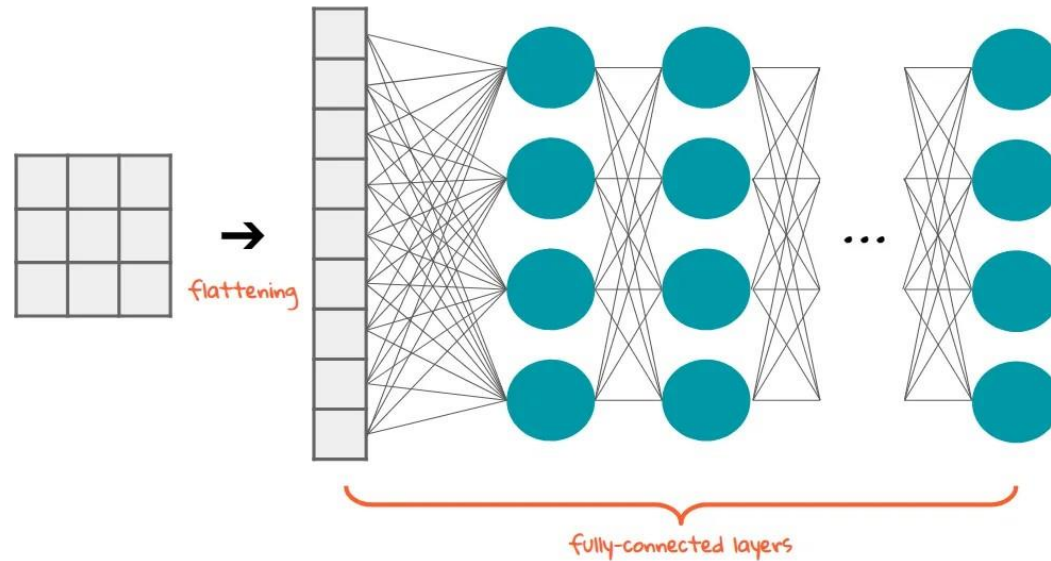


Flattening Layer

Converting multi-dimensional feature maps from the convolutional and pooling layers into a one-dimensional array.

Why? To connect extracted features to the final classification task. Fully-connected layers require linear input format.

How? Takes the output from the last pooling layer (dimensions: Height, Width, Depth) and reshapes into a single vector with a size Height x Width x Depth



Notion of Padding and Stride






Image: Jiwon Jeong, e2eml.schools

Padding: adding extra pixels around input data before convolution ('valid' when none, 'same' when output match input)

Stride: number of pixels a convolution filter moves/ steps at a time


Crucial for effective feature extraction (relevant features) and model performance

How to get a feature map?



7x7 Image (padded)

*



3x3 kernel

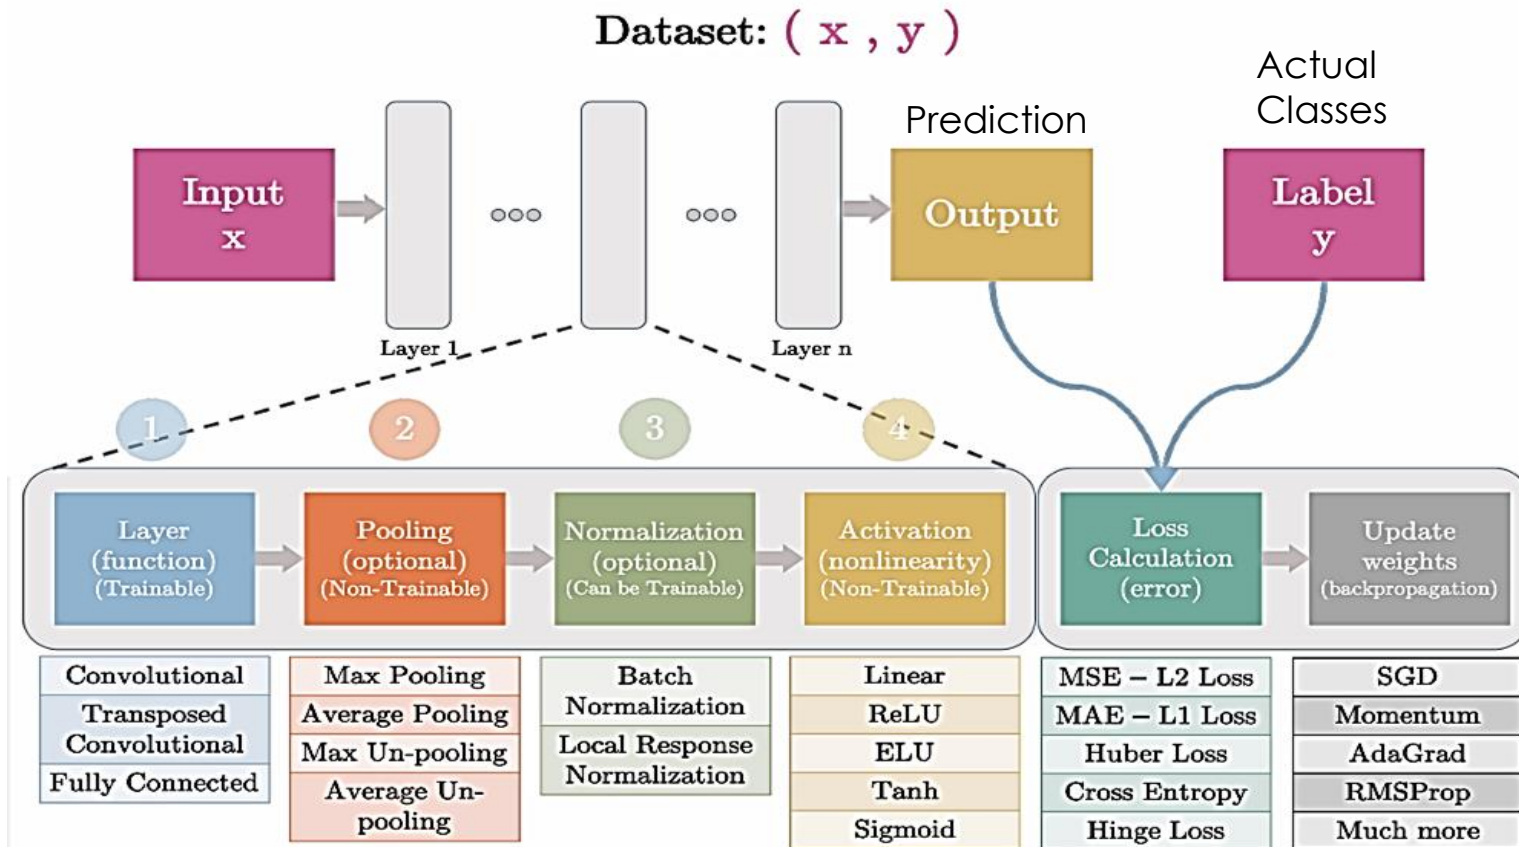
1	2	1	1	0

5x5 Feature Map

Training techniques

- **Relu** activation function introduces non-linearity
- **Dropout** regularization prevents overfitting
- **Data Augmentation** provides more data (perspectives, side views, 3D etc)
- **Transfer Learning** reduces time of training and resources ex. deriving other feline families once the concept of a cat is understood by the model

Trainable means
can be fine-tuned:
Hyperparameters
weights and biases



Tools for Computer Vision using CNNs

Frameworks

- **TensorFlow:** An open-source library for deep learning developed by Google, ideal for building CNNs.
- **PyTorch:** A flexible deep learning framework from Facebook, popular for research and production.
- **Keras:** A high-level API that simplifies building CNNs, often used with TensorFlow.
- **MATLAB:** Provides tools for designing and deploying CNNs through its Deep Learning Toolbox



CPU: Central Processing Unit
GPU: Graphic Processing Unit
TPU: Tensor Processing Unit




Filters


Original
input image



→ Horizontal edge filter




→ Vertical edge filter




STRIDES

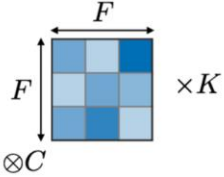
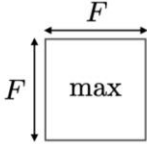
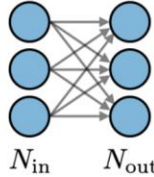
Horizontal edge filter with stride2



Vertical edge filter with stride2



Notion of complexity in training

	CONV	POOL	FC
Illustration			
Input size	$I \times I \times C$	$I \times I \times C$	N_{in}
Output size	$O \times O \times K$	$O \times O \times C$	N_{out}
Number of parameters	$(F \times F \times C + 1) \cdot K$	0	$(N_{\text{in}} + 1) \times N_{\text{out}}$
Remarks	<ul style="list-style-type: none">• One bias parameter per filter• In most cases, $S < F$• A common choice for K is $2C$	<ul style="list-style-type: none">• Pooling operation done channel-wise• In most cases, $S = F$	<ul style="list-style-type: none">• Input is flattened• One bias parameter per neuron• The number of FC neurons is free of structural constraints